

# Parallel Locally-Ordered Clustering for Bounding Volume Hierarchy Construction

Daniel Meister and Jiří Bittner

Department of Computer Graphics and Interaction  
Faculty of Electrical Engineering  
Czech Technical University in Prague



# Motivation: Interactive Ray Tracing

Fast BVH construction for geometry that is not known a priori



- Dynamic geometry changes in every frame
- Scene is assembled on the fly



[Benthin et al. 2017]

# Bounding Volume Hierarchy (BVH)

- Ray tracing, collision detection, visibility culling
- Rooted tree of arbitrary branching factor
  - References to geometric primitives in leaves
  - Bounding volumes in interior nodes




[Clark 1976]

# BVH Construction Methods




## Top-down

- Surface Area Heuristic [Hunt et al. 2007]
- Binning [Ize et al. 2007, Wald 2007]
- $k$ -means clustering [Meister and Bittner 2016]



## Bottom-up

- Agglomerative clustering [Walter et al. 2008]
- Approximate aggl. clustering [Gu et al. 2013]




# BVH Construction Methods




## Insertion

- Heuristic greedy search [Goldsmith and Salmon 1987]
- Online construction [Bittner et al. 2015]



## Optimization



- Rotations [Kensler 2008, Kopta et al. 2012]
- Insertion-based optimization [Bittner 2013 et al.]
- Treelet restructuring [Karras and Aila 2013, Domingues and Pedrini 2015]



# Surface Area Heuristic (SAH)



$$c(N) = \begin{cases} c_T + \frac{SA(N_L)}{SA(N)}c(N_L) + \frac{SA(N_R)}{SA(N)}c(N_R) & \text{if } N \text{ is interior node} \\ c_I|N| & \text{otherwise} \end{cases}$$





[MacDonald and Booth 1990]

# Surface Area Heuristic (SAH)



$$c(N) = \begin{cases} c_T + \frac{SA(N_L)}{SA(N)} c(N_L) + \frac{SA(N_R)}{SA(N)} c(N_R) & \text{if } N \text{ is interior node} \\ c_I |N| & \text{otherwise} \end{cases}$$





[MacDonald and Booth 1990]

# Surface Area Heuristic (SAH)



$$c(N) = \begin{cases} c_T + \frac{SA(N_L)}{SA(N)} c(N_L) + \frac{SA(N_R)}{SA(N)} c(N_R) & \text{if } N \text{ is interior node} \\ c_I |N| & \text{otherwise} \end{cases}$$



[MacDonald and Booth 1990]





# Surface Area Heuristic (SAH)



$$c(N) = \begin{cases} c_T + \frac{SA(N_L)}{SA(N)} c(N_L) + \frac{SA(N_R)}{SA(N)} c(N_R) & \text{if } N \text{ is interior node} \\ c_I |N| & \text{otherwise} \end{cases}$$

$$c(N_{root}) = \frac{1}{SA(N_{root})} \left[ c_T \sum_{N_i} SA(N_i) + c_I \sum_{N_l} SA(N_l) |N_l| \right]$$





[MacDonald and Booth 1990]

# Surface Area Heuristic (SAH)



$$c(N) = \begin{cases} c_T + \frac{SA(N_L)}{SA(N)} c(N_L) + \frac{SA(N_R)}{SA(N)} c(N_R) & \text{if } N \text{ is interior node} \\ c_I |N| & \text{otherwise} \end{cases}$$

$$c(N_{root}) = \frac{1}{SA(N_{root})} \left[ c_T \sum_{N_i} SA(N_i) + c_I \sum_{N_l} SA(N_l) |N_l| \right] \propto \sum_{N_i} SA(N_i) \text{ if } |N_l| = 1$$



[MacDonald and Booth 1990]

# Agglomerative Clustering



# Agglomerative Clustering

Repeat until only one cluster remains



# Agglomerative Clustering

Repeat until only one cluster remains

- Search for nearest neighbors for each cluster



# Agglomerative Clustering



Repeat until only one cluster remains

- Search for nearest neighbors for each cluster
- Merge the *closest* cluster pair

# Agglomerative Clustering




Repeat until only one cluster remains


- Search for nearest neighbors for each cluster
- Merge the *closest* cluster pair

Distance between clusters  $C_1$  and  $C_2$  [Walter et al. 2008]

$$d(C_1, C_2) = SA(C_1 \cup C_2)$$




# Agglomerative Clustering







# Agglomerative Clustering




# Agglomerative Clustering




# Agglomerative Clustering




# Agglomerative Clustering





# Agglomerative Clustering





# Agglomerative Clustering



# Agglomerative Clustering




# Agglomerative Clustering







# Agglomerative Clustering




# Agglomerative Clustering




# Agglomerative Clustering



# Agglomerative Clustering




# Agglomerative Clustering



# Locally-Ordered Clustering

Non-decreasing property [Walter et al. 2008]

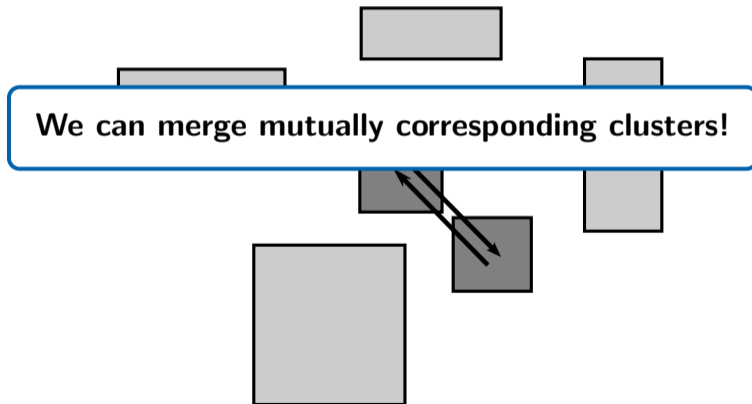
$$d(C_1, C_2) \leq d(C_1 \cup C_3, C_2) : \forall C_1, C_2, C_3$$




# Locally-Ordered Clustering

Non-decreasing property [Walter et al. 2008]

$$d(C_1, C_2) \leq d(C_1 \cup C_3, C_2) : \forall C_1, C_2, C_3$$




# Parallel Locally-Ordered Clustering

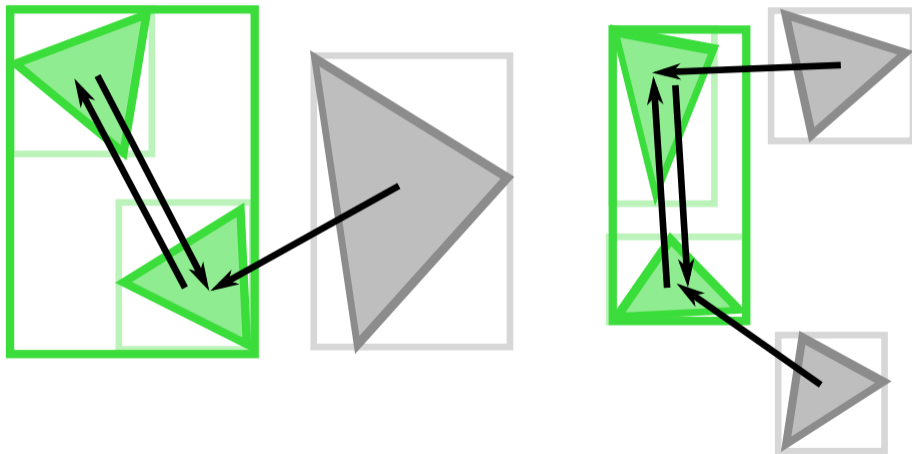





# Parallel Locally-Ordered Clustering




# Parallel Locally-Ordered Clustering




# Parallel Locally-Ordered Clustering





# Parallel Locally-Ordered Clustering




# Parallel Locally-Ordered Clustering




# Parallel Locally-Ordered Clustering



# Parallel Locally-Ordered Clustering




# Parallel Locally-Ordered Clustering





# Parallel Locally-Ordered Clustering



# Nearest Neighbor Search



# Nearest Neighbor Search



Naïve approach

- Time complexity  $\mathcal{O}(n^2)$
- Prohibitive for practical use

# Nearest Neighbor Search



Naïve approach

- Time complexity  $\mathcal{O}(n^2)$
- Prohibitive for practical use

kD-tree [Walter et al. 2008]

- Difficult to implement
- Not suitable for parallel processing

# Nearest Neighbor Search



Naïve approach

- Time complexity  $\mathcal{O}(n^2)$
- Prohibitive for practical use

kD-tree [Walter et al. 2008]

- Difficult to implement
- Not suitable for parallel processing

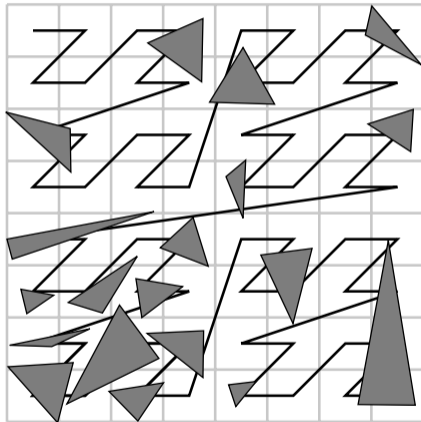
Morton curve (our approach)

- Sort clusters along the Morton curve
- Search in the sorted array around a given cluster
- Neighborhood around the cluster determined by radius parameter  $r$

# Approx. Nearest Neighbors along Morton Curve



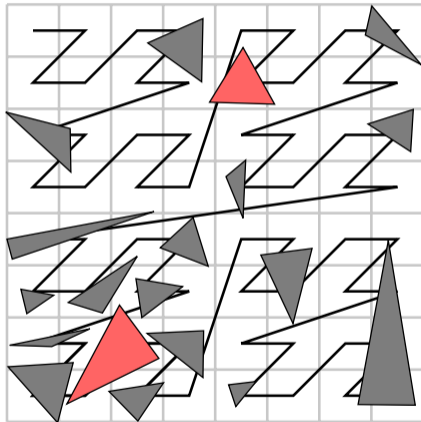
Neighborhood determined by radius  $r = 2$



# Approx. Nearest Neighbors along Morton Curve




Neighborhood determined by radius  $r = 2$



# Approx. Nearest Neighbors along Morton Curve



Neighborhood determined by radius  $r = 2$





# Algorithm Overview



# Algorithm Overview


Repeat until one cluster remains



# Algorithm Overview

Repeat until one cluster remains


- Search nearest neighbor (in parallel)



# Algorithm Overview

Repeat until one cluster remains

- Search nearest neighbor (in parallel)




- Merge (in parallel)



# Algorithm Overview

Repeat until one cluster remains


- Search nearest neighbor (in parallel)




- Merge (in parallel)



- Compact via prefix scan (in parallel)




# Parallel Subtree Collapsing



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off





# Parallel Subtree Collapsing


- 1 Decide whether collapsing pays off



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off



# Parallel Subtree Collapsing



- 1 Decide whether collapsing pays off






# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)



# Parallel Subtree Collapsing



- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)






# Parallel Subtree Collapsing



- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)






# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid



# Parallel Subtree Collapsing


- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid



# Parallel Subtree Collapsing




- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid



# Parallel Subtree Collapsing



- 1 Decide whether collapsing pays off
- 2 Identify leaf nodes (i.e. roots of collapsed subtrees)
- 3 Mark nodes as valid or invalid






# Implementation in CUDA

Shared memory cache of size  $B + 2r$


- Block with  $B$  threads
- Radius  $r$



# Results



- 9 scenes (331k - 12759k tris)
- Path tracing (GPU ray tracing kernel [Aila and Laine 2009])
  - Low quality rendering (8 spp)
  - High quality rendering (128 spp)
- Intel Core i7-3770 3.4 GHz CPU (4 cores), 16 GB RAM
- NVIDIA GeForce GTX TITAN X (Maxwell), 12 GB RAM



# Tested Methods



LBVH [Karras 2012]

- Spatial median splits

HLBVH [Garanzha et al. 2011]

- Spatial median and SAH splits

ATRBVH [Domingues and Pedrini 2015]

- Treelet restructuring by agglomerative clustering



PLOC

- Parallel locally-ordered clustering (our algorithm)

Adaptive leaf sizes, SAH cost constants  $c_T = 3$ ,  $c_I = 2$



# Pompeii

SAH cost (5632k tris,  $r = 25$ )





# Pompeii

build time (5632k tris,  $r = 25$ )





# Pompeii

time-to-image LQ (5632k tris,  $r = 25$ )





# Pompeii

time-to-image HQ (5632k tris,  $r = 25$ )



# Powerplant



SAH cost (12759k tris,  $r = 10$ )






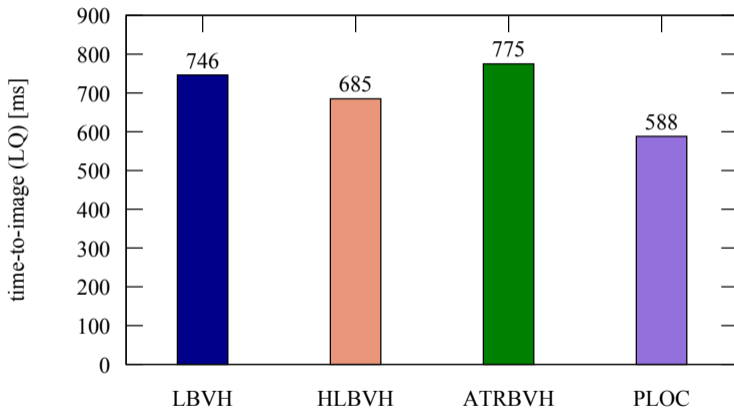
# Powerplant

build time (12759k tris,  $r = 10$ )




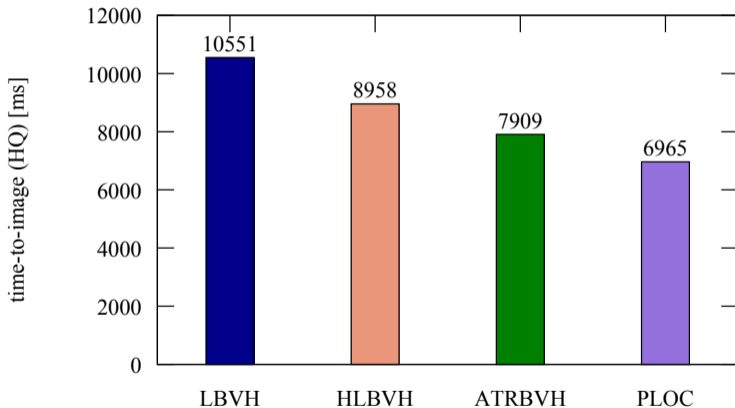
# Powerplant

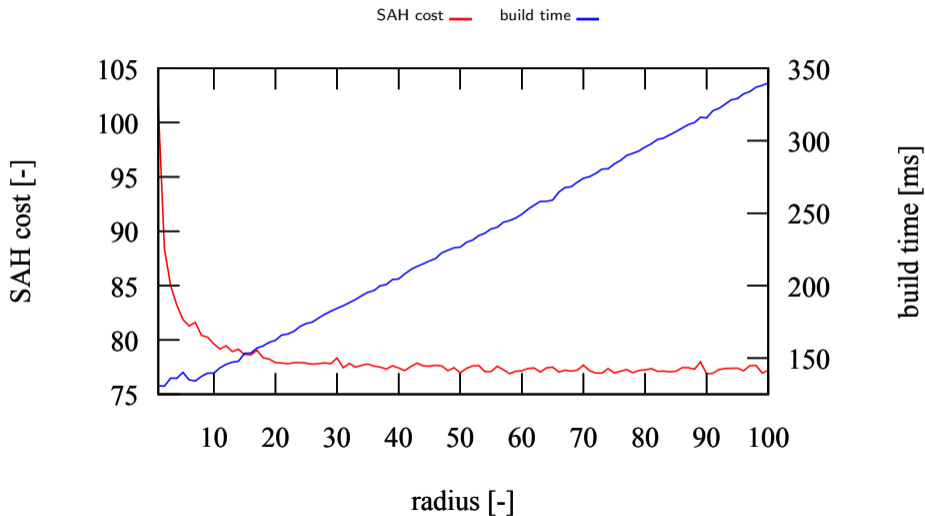
time-to-image LQ (12759k tris,  $r = 10$ )




# Powerplant

time-to-image HQ (2759k tris,  $r = 10$ )







# Conclusion and Future Work



GPU-based BVH construction using appr. agglomerative clustering

- Efficient and extremely simple
- Parallel subtree collapsing
- Implementation in CUDA with released source codes

Future work

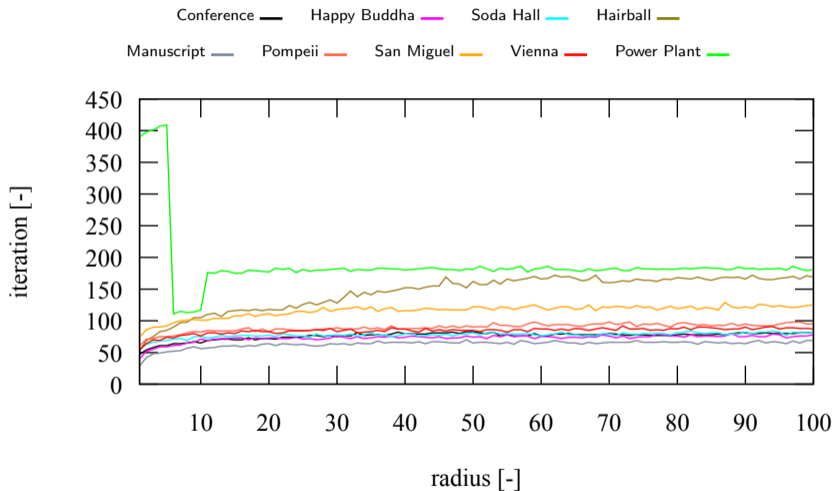
- Varying radius across different iterations
- Extended Morton codes [Vinkler et al. 2017]

# Thank you for your attention!

The project website with source codes  
<http://dcgi.fel.cvut.cz/projects/ploc/>



# Iterations





# Comparison with AAC

Approximate agglomerative clustering [Gu et al. 2013]

